

PEOPLE AND CULTURE GOVERNANCE

Recruitment and AI-Assisted Hiring Policy

Controls the approved use of AI systems and agents in candidate screening, ranking, recommendation, communication, and employment decision support.

ACTIVE**Version 1.4**

Synthetic demonstration document. Prepared for the PolicyForge prototype. It contains no real candidate, employee, or company data.

Document control

DOCUMENT ID POL-HR-014	VERSION 1.4
POLICY OWNER Director of Human Resources	ACCOUNTABLE EXECUTIVE Chief People Officer
APPROVAL AUTHORITY AI Governance Committee	STATUS Active
EFFECTIVE DATE 3 January 2026	NEXT REVIEW 3 January 2027

Contents

1. Purpose	3	2. Scope	3
3. Definitions	3	4. Candidate data	4
5. Approved AI use	4	6. Fair evaluation	4
7. Decisions and review	5	8. Candidate transparency	6
9. Evidence and audit	6	10. Exceptions and changes	6

1. Purpose

This policy ensures that AI-assisted recruitment protects candidate data, uses job-related evidence, preserves human responsibility, and provides candidates with understandable notice and a route to reassessment.

2. Scope

This policy applies to employees, contractors, recruitment partners, AI models, software agents, and connected systems used to source, screen, rank, recommend, communicate with, or make decisions affecting candidates for Northstar Group roles.

2.3

System boundary

The Northstar applicant tracking system remains the system of record. An AI agent may prepare analysis and recommendations only through approved integrations. A final employment outcome must remain attributable to an authorized human decision maker.

3. Definitions

TERM	
DEFINITION	
AI-assisted recruitment	Use of a model or software agent to process candidate information or support a recruitment action.
Candidate Screening Agent	The registered agent that reads approved application fields, requests approved model inference, ranks candidates, and drafts notices.
Consequential recommendation	A recommendation that may materially affect interview selection, ranking, rejection, or another employment opportunity.
External model	A model endpoint operated outside Northstar-controlled infrastructure.
Decision Capsule	An append-only reconstruction of recorded facts, model output, agent action, policy determination, human decision, and external consequence.

4. Candidate data

4.2 Identity minimization

Candidate name, contact details, address, and equivalent direct identifiers must be removed before application content is sent to an external model. The transformation and destination must be recorded.

4.4 Purpose limitation

Candidate information may be processed only for the approved recruitment purpose, role, and retention period recorded in the Agent Manifest.

5. Approved AI use

5.3 Approved models and tools

Resume content and employment history may be processed only by models and tools approved by Information Security and listed in the current Agent Manifest.

6. Fair evaluation

6.1 Protected attributes

Protected characteristics and reliable proxies must not be used to rank, profile, or disadvantage a candidate. Any attempted use must be blocked and recorded.

6.3 Job-related criteria

Candidate ranking may use documented qualifications, required certifications, relevant experience, and other approved job-related criteria. The criteria and policy version must be recorded.

Runtime outcomes. Approved controls must return one of ALLOW, TRANSFORM, HUMAN REVIEW, or BLOCK with the triggering clause and evidence reference.

7. Consequential decisions

7.1 Recommendation authority

An AI system may prepare a candidate ranking or recommendation. It must not issue a final rejection, selection, or employment outcome independently.

7.3 Human review

A negative candidate recommendation requires approval from an HR manager or assigned recruitment lead. The reviewer must receive the recorded factors, model and agent versions, applicable control, and policy version.

7.4 Employment history

Employment gaps may be considered when evaluating candidate suitability when the consideration is relevant to documented role requirements and remains subject to authorized human review.

7.6 Reviewer responsibility

The reviewer owns the final decision and must be identified in the evidence record. Approval of an agent recommendation does not transfer responsibility to the model or agent.

Control interpretation: a negative recommendation produced by Candidate Screening Agent v2.1 is paused before consequence and routed to an authorized reviewer under control CTL-025.

8. Candidate transparency

8.2 Notice and reassessment

A candidate notice must disclose material AI assistance, identify human responsibility for the outcome, summarize important recorded factors, and provide a route to human reassessment.

9. Evidence and audit

9.1 Decision evidence

Every consequential recommendation must create a Decision Capsule recording relevant inputs, transformations, model output, agent recommendation, policy determination, human decision, notice, and version provenance. Private model chain-of-thought must not be collected as audit evidence.

10. Exceptions and changes

Exceptions require written approval from the policy owner and AI Governance Committee. Changes must create a new immutable version and must not overwrite the policy applied to historical decisions.

VERSION			
DATE			
CHANGE			
APPROVAL			
1.4	3 Jan 2026	Added Agent Manifest, external-model transformation, Decision Capsule, and candidate reassessment requirements.	AI Governance Committee
1.3	9 Sep 2025	Added approved-model and protected-attribute controls.	Chief People Officer

Elena Park

Chief People Officer

Approved 3 January 2026

Marcus Lee

Chair, AI Governance Committee

Approved 3 January 2026